



SCALABLE GENOMIC COMPRESSION: TRANSFORMING DNA INTO BITS

Ryan Son, Gavin Sixer, Jared Arroyo Ruiz | Advised by Professor Layla Oesper
Department of Computer Science, Carleton College



Introduction

Genome sequencing has dropped from \$30 million to just **\$80 per genome** over the past two decades, removing financial barriers to large-scale research [1]. But with this affordability comes a **flood of data**—making **efficient compression algorithms essential** for cost-effective storage and analysis. We aimed to explore, implement, and test several genomic data algorithms: **DNAZip**, Biocompress, and Huffman Coding.

Compression Strategies

Type	Algorithm	Key Feature
Reference	DNAZip (2009)	Uses a reference genome to compare against a target genome and compress the -0.1% of differences [2]
Non-reference	Biocompress (1993)	Exploits intrinsic sequence properties like repetitions without a reference [3]
Entropy-based	Huffman Coding (1952)	Ordered assigning of the most frequent characters/symbols/words/k-mer with the shortest bitstring encodings [4]

Compressing DNA with DNAZip

DNA is composed of four **unique** nucleotides:

- Adenine (A)
- Guanine (G)
- Cytosine (C)
- Thymine (T)

Three types of possible **variations** may exist at any given position:




Figure 1. Target genome variations with respect to the reference genome: SNVs – single nucleotide variants; deletions – loss of 1+ nucleotides; and insertions – addition of 1+ nucleotides.

The **DNAZip** algorithm harnesses the fact that human genomes are **99.9% identical** [2]. This requires a **reference genome**, a **target genome**, and a **SNP database**. The target is **compared** against the reference to **create** a variant call format (VCF) file. This **VCF** file contains the **differences/variations** of the **target genome**.

Variation Header	Variation Info
SNV (0), CHROMOSOME, POSITION	C/G
DEL (1), CHROMOSOME, POSITION	TG--
INS (2), CHROMOSOME, POSITION	--GA

Table 1. Formatting of a VCF file for any target genome.




Figure 2. Creation of the VCF utilized **MUMmer4** [5] (genome alignment tool), a seed-extend algorithm that finds reference/target nucleotide matches and then dynamically identifies the DNA variations and **all2vcf** [6], a VCF conversion tool.

The processed VCF file is **compressed** by DNAZip through the combination of several key methods: **VINTs** (variable integer lengths), **SNP mapping**, **Huffman coding** of insertions, and **delta position encoding**.

Methodology




Figure 3. DNAZip encoding of a VCF consists of four key parts: VINTs, dbSNP, Huffman Coding, and Delta Positions. Decoding reconstructs the original data by reversing each part.

8-mer	Bitstring	Length
CTACATAT	00000000000000	15
CCTAAGGT	1111111100001111	17
GCAATCGC	11111111101110101	20
CGGTACAA	11111101101100100000	22

Figure 4. VINT representations only require the least number of bytes to store any integer.

Table 3. Sample k-mers ($k = 8$) from the Huffman encoding map. Each distinct k-mer has a unique bitstring. Each bitstring is encoded into the compressed file when the matching k-mer has been found. Ranked by frequency from top to bottom. The most frequent k-mer has the shortest bitstring.

DNAZip **encodes** all **integers** in a VCF—including variant positions, bitstring lengths, and deletion lengths—as **VINTs**. For each chromosome, mapped SNPs are stored in a bitstring, while unmapped SNVs are encoded using a two-bit nucleotide representation [2].

Insertions are compressed with **Huffman coding**: all insertion sequences of a chromosome are first concatenated, then partitioned into 8-mers, and finally frequency-analyzed to assign **shorter bitstrings** to the most **common** 8-mers. Bitstring lengths are recorded per variant type to support decoding.




Figure 5. Increasing k-mer sizes when Huffman coding on several genomes resulted in higher space savings, or more effective compression. 8-mers provided the greatest space savings, so a k-mer of size 8 has been chosen for DNAZip.

After encoding each variation type of each chromosome, the final **bitstring** is then **stored** in the final compressed file.

Decoding reads the compressed file into a bitstring and uses the stored bitstring lengths to determine how far to advance through each variant type and chromosome.

Each compression feature can be individually enabled or disabled to quantify its contribution to overall VCF size reduction.

Results




Figure 6. Per chromosome **construction** of the encoded bitstring in the compressed file containing each variant. Delta positions VINTs, variant length VINTs, and two-bit encodings are concatenated. The VINT (length) allows the decoder to intuitively know how many bits (positions, variant info) to iterate through.




Figure 7. Compressed file sizes depending on the activated compression features. dbSNP adversely contributed to compression by increasing the size of the end file. Huffman coding marginally improved compression. Delta positions provided the largest decrease in compressed file size.

Adding dbSNPs degraded compression performance. As this DNAZip implementation included only Common SNPs, very few SNVs in the target genomes actually mapped to entries in dbSNP. As a result, the bitstring grew without providing meaningful variation data, needlessly inflating the output.

In contrast, **delta-encoding** variant positions yielded the largest compression gains, **reducing** the encoded file **size** by nearly **50%**. This improvement stems from the large absolute genomic positions: encoding positional differences dramatically lowers the magnitude of these values, allowing the **VINT** representations to operate much more **efficiently** and occupy less space, greatly improving compression performance.




Figure 8. Overall compressed genome sizes across four types, Huffman coding (8-mer), Biocompress 1, uncompressed VCF, and DNAZip compression of the VCF. DNAZip has the smallest end file size, whereas Huffman has the largest.

Conclusion

Reference-based tools like DNAZip demonstrate exceptional scalability by efficiently representing genomic differences instead of entire sequences.




Figure 9. Linear growth of each algorithm. Starting storage sizes are equivalent to a single file encoded by each algorithm, with the exception of DNAZip including the uncompressed size of both its dbSNP and reference genome (~3.5 GB).

While DNAZip offers **superior space savings** compared to its non-reference-based counterparts like Biocompress. It also introduces additional computational complexity by **requiring both pre-processing** of the **target** genome (alignment for VCF creation) and storage of **reference** data.

Overall, referential compression systems like DNAZip represent a crucial advancement toward **scalable, cost-effective** genomic data management in the era of high-throughput **sequencing**.

Future Work

Test larger dbSNPs to determine whether higher mapping rates justify the storage cost of a substantially expanded SNP database and identify the number of human genomes where this tradeoff becomes beneficial. Additionally, more modern reference-based compressor would be explored and implemented to test their performance against DNAZip.

Acknowledgements

I would like to thank my amazing group members for their incredible contributions, Technical Director Mike Tie for his technical assistance, and Professor Layla Oesper for her invaluable guidance and support.

References

- [1] F. L. Genomics and L. Fletcher, "The \$100 Genome: Where's the Limit?", *Front Line Genomics*. Accessed: Oct. 05, 2025. [Online]. Available: <https://frontlinegenomics.com/the-100-genome-wheres-the-limit/>
- [2] S. Christley, Y. Lu, C. Li, and X. Xie, "Human genomes as email attachments," *Bioinformatics*, vol. 25, no. 2, pp. 274–275, Jan. 2009, doi: 10.1093/bioinformatics/btn582.
- [3] S. Grumbach and F. Tahi, "Compression of DNA sequences," in *Proc. Data Compression Conf.*, 1993, pp. 340–350, doi: 10.1109/dcc.1993.253115.
- [4] D. A. Huffman, "A method for the construction of minimum-redundancy codes," *Reson*, vol. 11, no. 2, pp. 91–99, Feb. 2006, doi: 10.1007/BF02837279.
- [5] G. Marçais, A. L. Delcher, A. M. Phillip, R. Coston, S. L. Salzberg, and A. Zimin, "MUMLer4: A fast and versatile genome alignment system," *PLoS Comput Biol*, vol. 14, no. 1, p. e1005944, Jan. 2018, doi: 10.1371/journal.pcbi.1005944.
- [6] M. Schiavino, "Matteo.schiavino/all2vcf: Toolkit to convert the output of common variant calling programs to VCF," GitHub, <https://github.com/Matteo.schiavino/all2vcf> (accessed Oct. 9, 2025).